

Predicting HOF Members

May 30, 2024

Author: Mitchell vom Scheidt

1 Question: How accurately can we predict MLB Hall of Fame members based on their offensive performance?

1.0.1 Type of Problem: Supervised, Classification Machine Learning

1.0.2 Steps involved in this classification machine learning problem:

1. Find/import the data
2. Data preprocessing
3. Exploratory data analysis (EDA)
4. Feature engineering & selection
5. Model training
6. Model evaluation
7. Conclusion

1.1 Find/Import The Data

Dataset Description:

The data used in this project was taken from baseball reference. Specifically, the dataset contains MLB position players who debuted after 1946 (In order to include non-white players) and played their final game before 2018. Every player in the dataset played in ≥ 750 games in their career (Since the hall of fame requires members to play at minimum 764 games). The dataset contains 43 features including a wide range of statistics such as batting average (BA), wins above replacement (WAR), home runs (HR), and so on.

Target Variable: 'HOF Status'

```
[1]: import numpy as np
import pandas as pd
pd.options.display.max_columns = 50

data = pd.read_csv('data/past_player_data.csv')
data.head()
```

```
[1]:
```

	Player	HOF Status	Suspended	Suspected Steroids	WAR	\
0	Wayne Tolleson	False	False	False	2.3	
1	Gary Disarcina	False	False	False	11.2	

2	Darren Lewis	False	False	False	10.5
3	Brian Hunter	False	False	False	7.4
4	Alvaro Espinoza	False	False	False	3.9

	First Season	Last Season	Debut	Age	Final Season	Age	G	PA	AB	\
0	1981	1990		25		34	863	2614	2322	
1	1989	2000		21		32	1086	4032	3744	
2	1990	2002		22		34	1354	4654	4081	
3	1994	2003		23		32	1000	3659	3347	
4	1984	1997		22		35	942	2659	2478	

	R	H	1B	2B	3B	HR	xbh	RBI	SB	CS	BB	SO	BA	OBP	\
0	301	559	473	60	17	9	86	133	108	41	219	384	0.241	0.307	
1	444	966	732	186	20	28	234	355	47	44	154	306	0.258	0.292	
2	607	1021	820	137	37	27	201	342	247	107	403	514	0.250	0.323	
3	500	882	683	146	28	25	199	241	260	61	243	581	0.264	0.313	
4	252	630	494	105	9	22	136	201	13	19	76	324	0.254	0.279	

	SLG	OPS	OPS+	TB	GIDP	HBP	SH	SF	IBB	WAA	oWAR	dWAR	Rbat	\
0	0.293	0.600	66	680	40	8	53	11	0	-6.1	3.3	2.9	-98	
1	0.341	0.633	66	1276	105	36	77	21	0	-2.9	5.2	12.8	-164	
2	0.322	0.645	73	1313	58	48	101	19	1	-4.5	5.2	6.2	-123	
3	0.346	0.660	72	1159	46	13	27	29	1	-4.4	5.1	3.2	-119	
4	0.331	0.610	66	819	66	16	73	16	1	-5.1	1.0	7.3	-123	

	Rdp	Rbaser	Rbaser + Rdp	Rfield
0	8	3	11	-7
1	2	-1	2	63
2	14	2	16	53
3	11	30	41	27
4	-1	0	-1	31

1.2 Data Preprocessing

```
[2]: from matplotlib import pyplot as plt
    %matplotlib inline
    import seaborn as sns
    from sklearn.model_selection import train_test_split, GridSearchCV
    from sklearn.preprocessing import StandardScaler
    from sklearn.metrics import accuracy_score, confusion_matrix
    from sklearn.linear_model import LogisticRegression
    from sklearn.neighbors import KNeighborsClassifier
    from sklearn.tree import DecisionTreeClassifier
    from sklearn.ensemble import RandomForestClassifier
```

Due to the real-world bias imposed on players who have been guilty or suspects of steroid use, we can remove those players from our dataset. After which, we can drop the 2 columns relating to this from the dataset.

```
[3]: data = data[(data['Suspended'] == False) & data['Suspected Steroids'] == False]
data = data.drop(columns=['Suspended', 'Suspected Steroids'])
```

1.3 Exploratory Data Analysis (EDA)

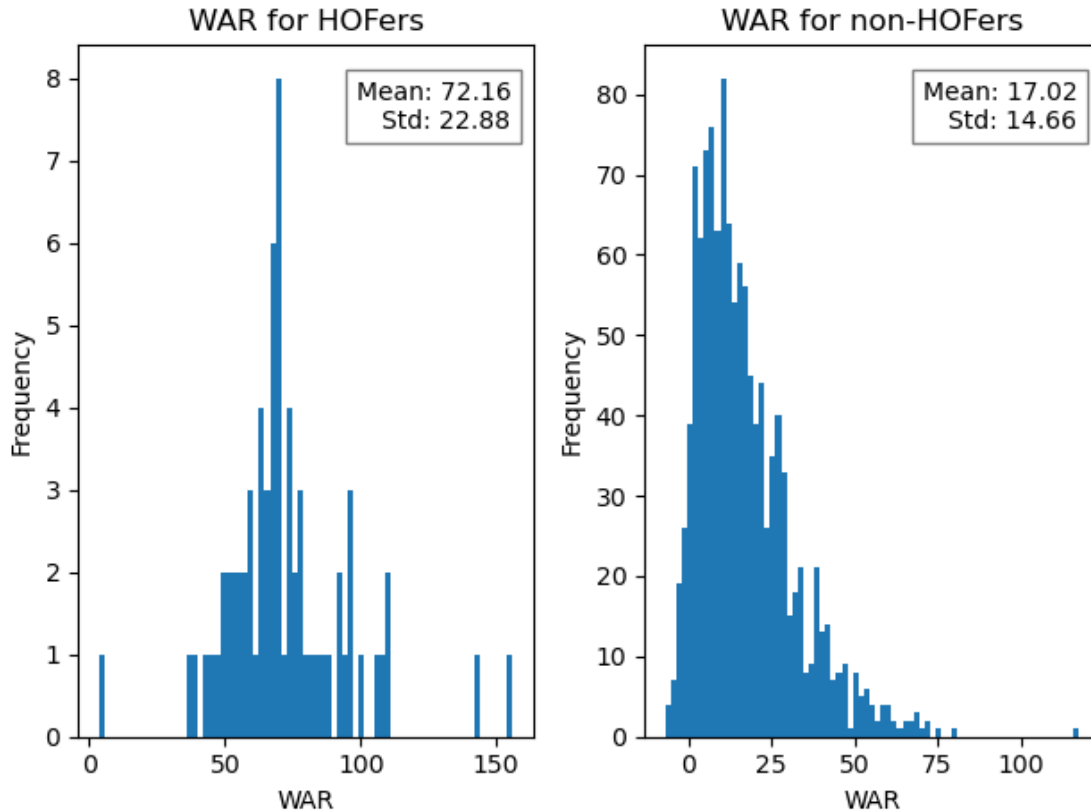
Let's explore how the values of some of the features in our dataset differ between players who are members of the hall of fame, and those which are not.

```
[4]: # helper function to creating histogram plot
def plotHist(var, bins, title, x_label, y_label):
    plt.hist(var, bins=bins)
    plt.title(title)
    plt.ylabel(y_label)
    plt.xlabel(x_label)
```

```
[5]: plt.subplot(1,2,1)
hof_war = data.loc[data['HOF Status'] == True, 'WAR']
# Creating plot
plotHist(hof_war, 75, 'WAR for HOFers', 'WAR', 'Frequency')
mean = np.mean(hof_war)
std = np.std(hof_war)
# Adding mean/std to plot
plt.text(0.95, 0.95, f'Mean: {mean:.2f}\nStd: {std:.2f}',
        transform=plt.gca().transAxes,
        verticalalignment='top',
        horizontalalignment='right',
        bbox=dict(facecolor='white', alpha=0.5))

plt.subplot(1,2,2)
non_hof_war = data.loc[data['HOF Status'] == False, 'WAR']
# Creating plot
plotHist(non_hof_war, 75, 'WAR for non-HOFers', 'WAR', 'Frequency')
mean = np.mean(non_hof_war)
std = np.std(non_hof_war)
# Adding mean/std to plot
plt.text(0.95, 0.95, f'Mean: {mean:.2f}\nStd: {std:.2f}',
        transform=plt.gca().transAxes,
        verticalalignment='top',
        horizontalalignment='right',
        bbox=dict(facecolor='white', alpha=0.5))

plt.tight_layout()
plt.show()
```

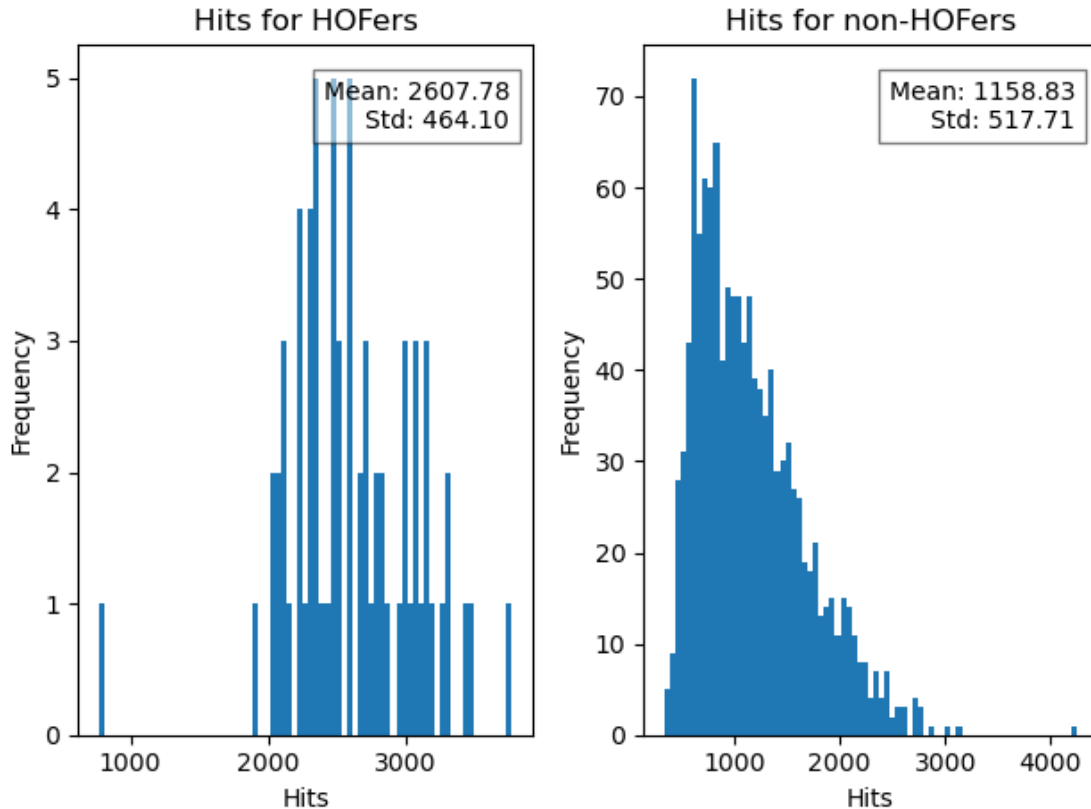


```
[6]: plt.subplot(1,2,1)
hof_hits = data.loc[data['HOF Status'] == True, 'H']
# Creating plot
plotHist(hof_hits, 75, 'Hits for HOFers', 'Hits', 'Frequency')
mean = np.mean(hof_hits)
std = np.std(hof_hits)
# Adding mean/std to plot
plt.text(0.95, 0.95, f'Mean: {mean:.2f}\nStd: {std:.2f}',
         transform=plt.gca().transAxes,
         verticalalignment='top',
         horizontalalignment='right',
         bbox=dict(facecolor='white', alpha=0.5))

plt.subplot(1,2,2)
non_hof_hits = data.loc[data['HOF Status'] == False, 'H']
# Creating plot
plotHist(non_hof_hits, 75, 'Hits for non-HOFers', 'Hits', 'Frequency')
mean = np.mean(non_hof_hits)
std = np.std(non_hof_hits)
# Adding mean/std to plot
```

```
plt.text(0.95, 0.95, f'Mean: {mean:.2f}\nStd: {std:.2f}',
        transform=plt.gca().transAxes,
        verticalalignment='top',
        horizontalalignment='right',
        bbox=dict(facecolor='white', alpha=0.5))

plt.tight_layout()
plt.show()
```



```
[7]: plt.subplot(1,2,1)
hof_avg = data.loc[data['HOF Status'] == True, 'BA']
# Creating plot
plotHist(hof_avg, 50, 'Batting Average for HOFers', 'Batting Average',
        ↪ 'Frequency')
mean = np.mean(hof_avg)
std = np.std(hof_avg)
# Adding mean/std to plot
plt.text(0.95, 0.95, f'Mean: {mean:.2f}\nStd: {std:.2f}',
        transform=plt.gca().transAxes,
        verticalalignment='top',
```

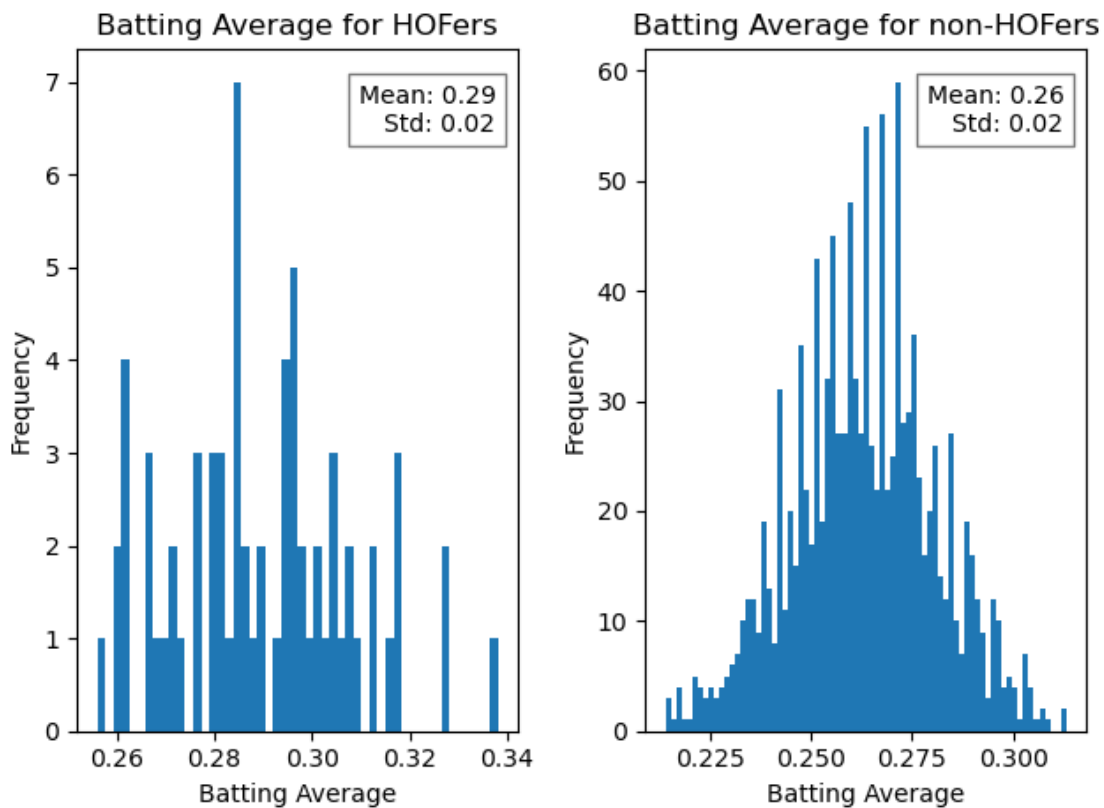
```

        horizontalalignment='right',
        bbox=dict(facecolor='white', alpha=0.5))

plt.subplot(1,2,2)
non_hof_avg = data.loc[data['HOF Status'] == False, 'BA']
# Creating plot
plotHist(non_hof_avg, 75, 'Batting Average for non-HOFers', 'Batting Average',
        ↪ 'Frequency')
mean = np.mean(non_hof_avg)
std = np.std(non_hof_avg)
# Adding mean/std to plot
plt.text(0.95, 0.95, f'Mean: {mean:.2f}\nStd: {std:.2f}',
        transform=plt.gca().transAxes,
        verticalalignment='top',
        horizontalalignment='right',
        bbox=dict(facecolor='white', alpha=0.5))

plt.tight_layout()
plt.show()

```



As you can see, for those statistics which are cumulative over a player's entire career, such as WAR

and Hits, the average and distribution of the values are drastically different between players which are members of the hall of fame and players which are not. In other words, these plots demonstrate the degree of separations between the two categories of players.

This will become important when we get to the feature engineering and selection process.

1.4 Feature Engineering & Selection

To start off, I have selected a subset of features which I will use and have discarded the rest of the features.

```
[8]: features = ['R', 'H', '1B', '2B', '3B', 'HR', 'xbh', 'RBI', 'SB', 'CS', 'BB', 'IBB', 'SO', 'TB', 'WAR', 'WAA']
other_cols = ['Player', 'HOF Status', 'G', 'BA', 'OBP', 'SLG', 'OPS']

data = data[other_cols + features]

data.head()
```

```
[8]:
```

	Player	HOF Status	G	BA	OBP	SLG	OPS	R	H	\
0	Wayne Tolleson	False	863	0.241	0.307	0.293	0.600	301	559	
1	Gary Disarcina	False	1086	0.258	0.292	0.341	0.633	444	966	
2	Darren Lewis	False	1354	0.250	0.323	0.322	0.645	607	1021	
3	Brian Hunter	False	1000	0.264	0.313	0.346	0.660	500	882	
4	Alvaro Espinoza	False	942	0.254	0.279	0.331	0.610	252	630	

	1B	2B	3B	HR	xbh	RBI	SB	CS	BB	IBB	SO	TB	WAR	WAA
0	473	60	17	9	86	133	108	41	219	0	384	680	2.3	-6.1
1	732	186	20	28	234	355	47	44	154	0	306	1276	11.2	-2.9
2	820	137	37	27	201	342	247	107	403	1	514	1313	10.5	-4.5
3	683	146	28	25	199	241	260	61	243	1	581	1159	7.4	-4.4
4	494	105	9	22	136	201	13	19	76	1	324	819	3.9	-5.1

Looking back to the EDA process, the feature values are different for the two classes of player (hall of fame & non-hall of fame). To eliminate this effect, I engineered features which average the features over a 162 game season.

```
[9]: for feature in features:
      data[f'{feature}_per_season'] = (data[feature] / data['G']) * 162

data.head()
```

```
[9]:
```

	Player	HOF Status	G	BA	OBP	SLG	OPS	R	H	\
0	Wayne Tolleson	False	863	0.241	0.307	0.293	0.600	301	559	
1	Gary Disarcina	False	1086	0.258	0.292	0.341	0.633	444	966	
2	Darren Lewis	False	1354	0.250	0.323	0.322	0.645	607	1021	
3	Brian Hunter	False	1000	0.264	0.313	0.346	0.660	500	882	
4	Alvaro Espinoza	False	942	0.254	0.279	0.331	0.610	252	630	

	1B	2B	3B	HR	xbh	RBI	SB	CS	BB	IBB	SO	TB	WAR	WAA	\
0	473	60	17	9	86	133	108	41	219	0	384	680	2.3	-6.1	
1	732	186	20	28	234	355	47	44	154	0	306	1276	11.2	-2.9	
2	820	137	37	27	201	342	247	107	403	1	514	1313	10.5	-4.5	
3	683	146	28	25	199	241	260	61	243	1	581	1159	7.4	-4.4	
4	494	105	9	22	136	201	13	19	76	1	324	819	3.9	-5.1	

	R_per_season	H_per_season	1B_per_season	2B_per_season	3B_per_season	\
0	56.502897	104.933951	88.790267	11.263036	3.191194	
1	66.232044	144.099448	109.193370	27.745856	2.983425	
2	72.624815	122.158050	98.109306	16.391433	4.426883	
3	81.000000	142.884000	110.646000	23.652000	4.536000	
4	43.337580	108.343949	84.955414	18.057325	1.547771	

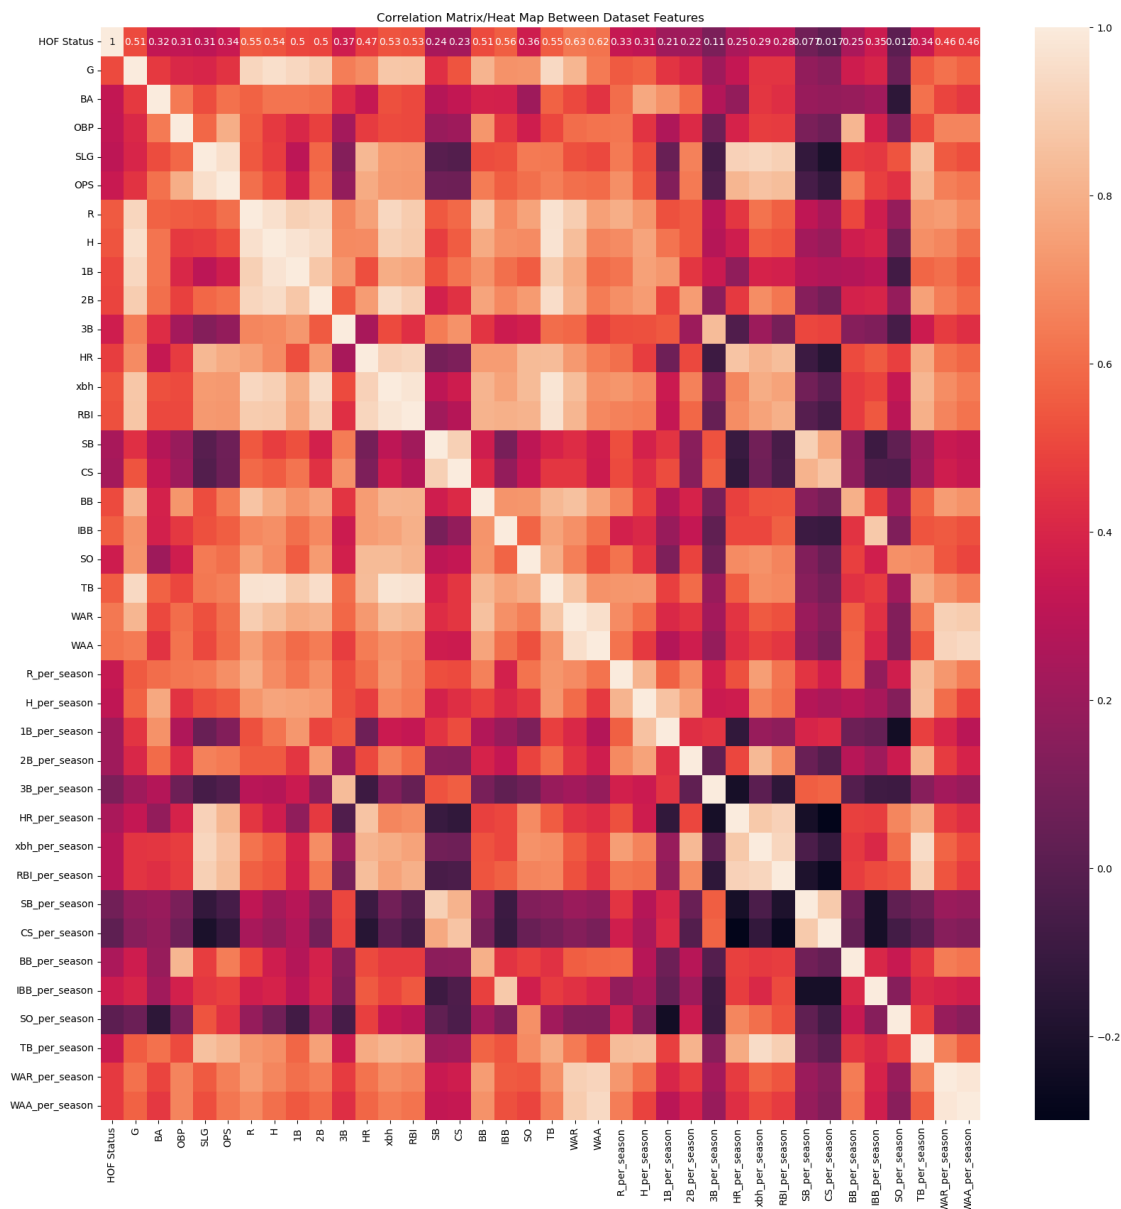
	HR_per_season	xbh_per_season	RBI_per_season	SB_per_season	\
0	1.689455	16.143685	24.966396	20.273465	
1	4.176796	34.906077	52.955801	7.011050	
2	3.230428	24.048744	40.918759	29.552437	
3	4.050000	32.238000	39.042000	42.120000	
4	3.783439	23.388535	34.566879	2.235669	

	CS_per_season	BB_per_season	IBB_per_season	SO_per_season	TB_per_season	\
0	7.696408	41.110081	0.000000	72.083430	127.647740	
1	6.563536	22.972376	0.000000	45.646409	190.342541	
2	12.802068	48.217134	0.119645	61.497784	157.094535	
3	9.882000	39.366000	0.162000	94.122000	187.758000	
4	3.267516	13.070064	0.171975	55.719745	140.847134	

	WAR_per_season	WAA_per_season
0	0.431750	-1.145075
1	1.670718	-0.432597
2	1.256278	-0.538405
3	1.198800	-0.712800
4	0.670701	-0.877070

Looking at a correlation matrix now using a heatmap we can see which features look promising and important to keep.

```
[10]: # Creating correlation matrix heatmap
plt.figure(figsize=(20,20))
sns.heatmap(data.corr(numeric_only=True), annot=True)
plt.title('Correlation Matrix/Heat Map Between Dataset Features')
plt.show()
```

```
[11]: data = data.drop(columns=features + ['G'])
      data.head()
```

```
[11]:
```

	Player	HOF Status	BA	OBP	SLG	OPS	R_per_season	\
0	Wayne Tolleson	False	0.241	0.307	0.293	0.600	56.502897	
1	Gary Disarcina	False	0.258	0.292	0.341	0.633	66.232044	
2	Darren Lewis	False	0.250	0.323	0.322	0.645	72.624815	
3	Brian Hunter	False	0.264	0.313	0.346	0.660	81.000000	
4	Alvaro Espinoza	False	0.254	0.279	0.331	0.610	43.337580	

	H_per_season	1B_per_season	2B_per_season	3B_per_season	HR_per_season	\
0	104.933951	88.790267	11.263036	3.191194	1.689455	
1	144.099448	109.193370	27.745856	2.983425	4.176796	
2	122.158050	98.109306	16.391433	4.426883	3.230428	
3	142.884000	110.646000	23.652000	4.536000	4.050000	
4	108.343949	84.955414	18.057325	1.547771	3.783439	

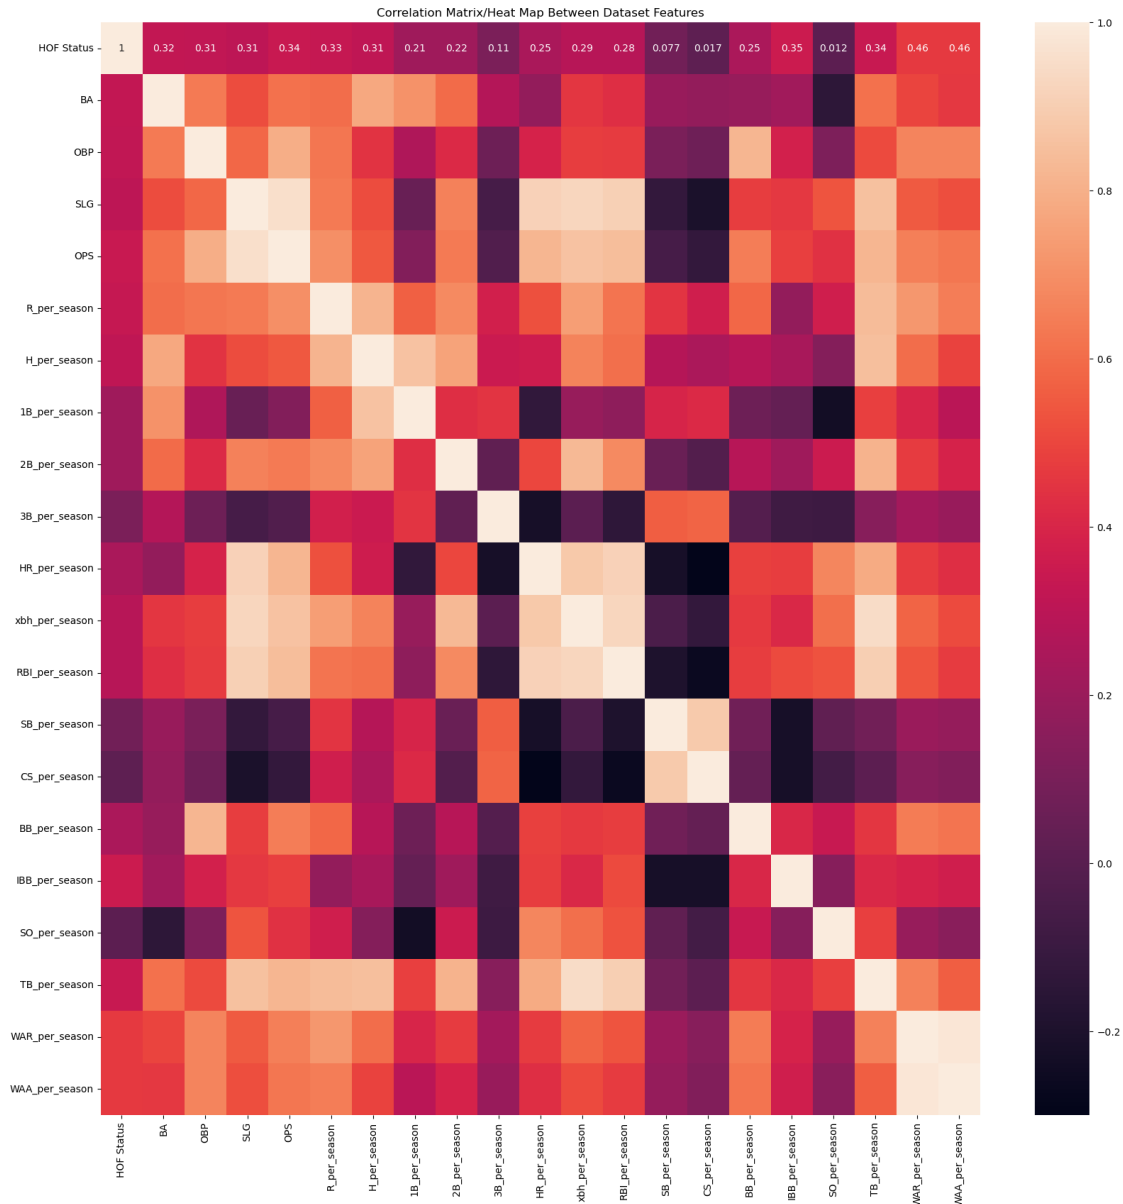
	xbh_per_season	RBI_per_season	SB_per_season	CS_per_season	\
0	16.143685	24.966396	20.273465	7.696408	
1	34.906077	52.955801	7.011050	6.563536	
2	24.048744	40.918759	29.552437	12.802068	
3	32.238000	39.042000	42.120000	9.882000	
4	23.388535	34.566879	2.235669	3.267516	

	BB_per_season	IBB_per_season	SO_per_season	TB_per_season	\
0	41.110081	0.000000	72.083430	127.647740	
1	22.972376	0.000000	45.646409	190.342541	
2	48.217134	0.119645	61.497784	157.094535	
3	39.366000	0.162000	94.122000	187.758000	
4	13.070064	0.171975	55.719745	140.847134	

	WAR_per_season	WAA_per_season
0	0.431750	-1.145075
1	1.670718	-0.432597
2	1.256278	-0.538405
3	1.198800	-0.712800
4	0.670701	-0.877070

Since we averaged the features in the engineering process on a per season basis, we can remove all the aggregate features and look at the correlation matrix heatmap again.

```
[12]: # Creating correlation matrix heatmap
plt.figure(figsize=(20,20))
sns.heatmap(data.corr(numeric_only=True), annot=True)
plt.title('Correlation Matrix/Heat Map Between Dataset Features')
plt.show()
```



Lastly, we need to change the 'HOF Status' column of our dataset, which is our target variable, to type integer instead of it being boolean.

```
[13]: # Converting boolean column to type integer
data['HOF Status'] = data['HOF Status'].astype(int)

final_data = data
final_data.head()
```

```
[13]:      Player  HOF Status  BA  OBP  SLG  OPS  R_per_season  \
0  Wayne Tolleson         0  0.241  0.307  0.293  0.600  56.502897
```

1	Gary Disarcina	0	0.258	0.292	0.341	0.633	66.232044
2	Darren Lewis	0	0.250	0.323	0.322	0.645	72.624815
3	Brian Hunter	0	0.264	0.313	0.346	0.660	81.000000
4	Alvaro Espinoza	0	0.254	0.279	0.331	0.610	43.337580

	H_per_season	1B_per_season	2B_per_season	3B_per_season	HR_per_season	\
0	104.933951	88.790267	11.263036	3.191194	1.689455	
1	144.099448	109.193370	27.745856	2.983425	4.176796	
2	122.158050	98.109306	16.391433	4.426883	3.230428	
3	142.884000	110.646000	23.652000	4.536000	4.050000	
4	108.343949	84.955414	18.057325	1.547771	3.783439	

	xbh_per_season	RBI_per_season	SB_per_season	CS_per_season	\
0	16.143685	24.966396	20.273465	7.696408	
1	34.906077	52.955801	7.011050	6.563536	
2	24.048744	40.918759	29.552437	12.802068	
3	32.238000	39.042000	42.120000	9.882000	
4	23.388535	34.566879	2.235669	3.267516	

	BB_per_season	IBB_per_season	SO_per_season	TB_per_season	\
0	41.110081	0.000000	72.083430	127.647740	
1	22.972376	0.000000	45.646409	190.342541	
2	48.217134	0.119645	61.497784	157.094535	
3	39.366000	0.162000	94.122000	187.758000	
4	13.070064	0.171975	55.719745	140.847134	

	WAR_per_season	WAA_per_season
0	0.431750	-1.145075
1	1.670718	-0.432597
2	1.256278	-0.538405
3	1.198800	-0.712800
4	0.670701	-0.877070

1.5 Model Training

As this is a classification problem, we need to use the appropriate machine learning models. The model's we will compare in this analysis are: 1. Logistic regression 2. K-nearest neighbors classifier 3. Decision tree classifier 4. Random forest classifier

Something important to note in the training process is stratification of the training data. We need to stratify our train/test datasets over the target variable since there are significantly more data entries which are non-hall of fame player than hall-of-fame players. Statifying the data means the proportion of data entries which are non-hall of fame to hall-of-fame is the same for both the train and test datasets.

I am using a 80/20 train/test split.

```
[14]: X = final_data.drop(columns=['Player', 'HOF Status'])
      y = final_data['HOF Status']

      # Creating train/test split
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪stratify=y, random_state=42)
```

We need to scale the data in order to standardize the features. We use the `sklearn.preprocessing.StandardScaler` which uses z-score standardization ($z = (x - u) / s$). It is necessary to standardize the features to ensure each feature is equally weighted. It also helps models such as logistic regression and k-nearest neighbors to converge faster and perform better.

```
[15]: # Scaling the data
      scaler = StandardScaler()
      X_train_scaled = scaler.fit_transform(X_train)
      X_test_scaled = scaler.transform(X_test)
```

When training some of the models, there is an opportunity for tuning/optimizing the hyperparameters. In this analysis, I chose to use `GridSearchCV` (grid search cross-validation), to optimize the hyperparameters for K-NN, Decision Tree, and Random Forest.

Logistic Regression:

```
[16]: # Creating & training model
      reg_model = LogisticRegression(max_iter=10000)
      reg_model.fit(X_train_scaled, y_train)
```

```
[16]: LogisticRegression(max_iter=10000)
```

K-Nearest Neighbors Classifier:

```
[17]: # Tuning hyperparameter n_neighbors (k)
      knn_model = KNeighborsClassifier()

      param_grid = {'n_neighbors': [1,3,5,7,8,11,13,15]}

      grid_search = GridSearchCV(estimator=knn_model, param_grid=param_grid, cv=10,
      ↪scoring='accuracy')
      grid_search.fit(X_train_scaled, y_train)

      optimal_n_neighbors = grid_search.best_params_['n_neighbors']
      print(f'Optimal number of neighbors: {optimal_n_neighbors}')

      # Creating & training model
      optimal_knn_model = KNeighborsClassifier(n_neighbors=optimal_n_neighbors)
      optimal_knn_model.fit(X_train_scaled, y_train)
```

Optimal number of neighbors: 15

```
[17]: KNeighborsClassifier(n_neighbors=15)
```

Decision Tree Classifier:

```
[18]: # Tuning hyperparameter max_depth
tree_model = DecisionTreeClassifier()

param_grid = {'max_depth': range(1, 20)}

grid_search = GridSearchCV(tree_model, param_grid, cv=10, scoring='accuracy')
grid_search.fit(X_train_scaled, y_train)

optimal_max_depth = grid_search.best_params_['max_depth']
print(f'Optimal max depth: {optimal_max_depth}')

# Creating & training model
optimal_tree_model = DecisionTreeClassifier(max_depth=optimal_max_depth,
    ↪random_state=42)
optimal_tree_model.fit(X_train_scaled, y_train)
```

Optimal max depth: 1

```
[18]: DecisionTreeClassifier(max_depth=1, random_state=42)
```

Random Forest Classifier:

```
[19]: # Tuning hyperparameter n_estimators
rf_model = RandomForestClassifier()

param_grid = {'n_estimators': [50, 100, 150, 200]}

grid_search = GridSearchCV(rf_model, param_grid, cv=10, scoring='accuracy')
grid_search.fit(X_train_scaled, y_train)

optimal_n_estimators = grid_search.best_params_['n_estimators']
print(f'Optimal n estimators: {optimal_n_estimators}')

# Creating & training model
optimal_rf_model = RandomForestClassifier(n_estimators=optimal_n_estimators,
    ↪random_state=42)
optimal_rf_model.fit(X_train_scaled, y_train)
```

Optimal n estimators: 200

```
[19]: RandomForestClassifier(n_estimators=200, random_state=42)
```

1.6 Model Evaluation

For classification problems, there are many ways to evaluate the performance of a model. The ones I will use are: 1. Accuracy 2. Precision 3. Recall 4. F1-Score

```

[20]: # Logistic Regression
y_pred_reg = reg_model.predict(X_test_scaled)

accuracy_reg = accuracy_score(y_test, y_pred_reg) * 100
conf_matrix_reg = confusion_matrix(y_test, y_pred_reg)
precision_reg = conf_matrix_reg[0,0] / (conf_matrix_reg[0,0] +
    ↪conf_matrix_reg[0,1]) * 100
recall_reg = conf_matrix_reg[0,0] / (conf_matrix_reg[0,0] +
    ↪conf_matrix_reg[1,0]) * 100
f1_reg = 2 * ((precision_reg * recall_reg) / (precision_reg + recall_reg)) / 100

# K-Nearest Neighbors Classifier
y_pred_knn = optimal_knn_model.predict(X_test_scaled)

accuracy_knn = accuracy_score(y_test, y_pred_knn) * 100
conf_matrix_knn = confusion_matrix(y_test, y_pred_knn)
precision_knn = conf_matrix_knn[0,0] / (conf_matrix_knn[0,0] +
    ↪conf_matrix_knn[0,1]) * 100
recall_knn = conf_matrix_knn[0,0] / (conf_matrix_knn[0,0] +
    ↪conf_matrix_knn[1,0]) * 100
f1_knn = 2 * ((precision_knn * recall_knn) / (precision_knn + recall_knn)) / 100

# Decision Tree Classifier
y_pred_tree = optimal_tree_model.predict(X_test_scaled)

accuracy_tree = accuracy_score(y_test, y_pred_tree) * 100
conf_matrix_tree = confusion_matrix(y_test, y_pred_tree)
precision_tree = conf_matrix_tree[0,0] / (conf_matrix_tree[0,0] +
    ↪conf_matrix_tree[0,1]) * 100
recall_tree = conf_matrix_tree[0,0] / (conf_matrix_tree[0,0] +
    ↪conf_matrix_tree[1,0]) * 100
f1_tree = 2 * ((precision_tree * recall_tree) / (precision_tree + recall_tree))
    ↪/ 100

# Random Forest Classifier
y_pred_rf = optimal_rf_model.predict(X_test_scaled)

accuracy_rf = accuracy_score(y_test, y_pred_rf) * 100
conf_matrix_rf = confusion_matrix(y_test, y_pred_rf)
precision_rf = conf_matrix_rf[0,0] / (conf_matrix_rf[0,0] +
    ↪conf_matrix_rf[0,1]) * 100
recall_rf = conf_matrix_rf[0,0] / (conf_matrix_rf[0,0] + conf_matrix_rf[1,0]) *
    ↪100
f1_rf = 2 * ((precision_rf * recall_rf) / (precision_rf + recall_rf)) / 100

```

```
# Creating a dataframe to display the results
results = pd.DataFrame({'Model':
                        ['Logistic Regression', 'K-Nearest Neighbors',
                        'Decision Tree Classifier', 'Random Forest Classifier'],
                        'Accuracy (%)':
                        [accuracy_reg , accuracy_knn, accuracy_tree,
                        accuracy_rf],
                        'Precision (%)':
                        [precision_reg, precision_knn, precision_tree,
                        precision_rf],
                        'Recall (%)':
                        [recall_reg, recall_knn, recall_tree, recall_rf],
                        'F1 Score':
                        [f1_reg, f1_knn, f1_tree, f1_rf]
                        })

results
```

```
[20]:
```

	Model	Accuracy (%)	Precision (%)	Recall (%)	\
0	Logistic Regression	94.531250	97.520661	96.721311	
1	K-Nearest Neighbors Classifier	95.312500	99.586777	95.634921	
2	Decision Tree Classifier	92.578125	95.041322	97.046414	
3	Random Forest Classifier	94.921875	98.347107	96.356275	

	F1 Score
0	0.971193
1	0.975709
2	0.960334
3	0.973415

1.7 Conclusion

The table above containing the various evaluation tools gives us a lot of insight. From it, we can tell: 1. Best overall model: K-Nearest Neighbors Classifier - Highest accuracy, precision, and F1 score 2. Highest Precision: K-Nearest Neighbors Classifier - Best for applications of minimizing false positives is needed 3. Highest Recall: Decision Tree Classifier - Best for applications where maximizing true positives is needed 4. Highest F1 Score: K-Nearest Neighbors Classifier - Best for applications where balancing precision and recall is needed

Overall, KNN came out to be the most effective model. Meaning it generalizes to new data well, and it is likely to perform reliably across different datasets. This project is a demonstration of key aspects of machine learning in general, however the model created could be used for, or tweaked slightly to then be used for, player analytics within baseball. That is, the evaluation of players to determine predicted future performance, and the suggested value of a given player.